

Predicting Houseprice:

This project was completed has part of Kaggle competition.

In this project we will dive deep into the following areas

1. Exploratory Data Analysis
 - a. Dependent variable
 - b. Univariate Analysis on Independent Variables
 - a. Dependent variable analysis
 - b. Univariate Analysis
 - A. Distribution of Numerical variables:
 - B. Distribution of Categorical Variables
 - c. Bivariate Analysis
 - C. Correlation among Numerical Variables
 2. Mean Disparity among Categorical Variables
 - d. Missing Values
 - e. Outlier detection

```
In [1]: ┌─ import pandas as pd
    import numpy as np
    import matplotlib.pyplot as plt
    import seaborn as sns
    from scipy import stats
    from scipy.stats import skew, norm
    from scipy.special import boxcox1p
    from scipy.stats import boxcox_normmax
    from scipy.stats.stats import pearsonr

    # Models
    from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
    from sklearn.kernel_ridge import KernelRidge
    from sklearn.linear_model import Ridge, RidgeCV
    from sklearn.linear_model import ElasticNet, ElasticNetCV
    from sklearn.svm import SVR
    from mlxtend.regressor import StackingCVRegressor
    import lightgbm as lgb
    from lightgbm import LGBMRegressor
    from xgboost import XGBRegressor

    # Stats
    from scipy.stats import skew, norm
    from scipy.special import boxcox1p
    from scipy.stats import boxcox_normmax

    # Misc
    from sklearn.model_selection import GridSearchCV
    from sklearn.model_selection import KFold, cross_val_score
    from sklearn.metrics import mean_squared_error
    from sklearn.preprocessing import OneHotEncoder
    from sklearn.preprocessing import LabelEncoder
    from sklearn.pipeline import make_pipeline
    from sklearn.preprocessing import scale
    from sklearn.preprocessing import StandardScaler
    from sklearn.preprocessing import RobustScaler
    from sklearn.decomposition import PCA

    # Ignore useless warnings
    import warnings
    warnings.filterwarnings(action="ignore")
    pd.options.display.max_seq_items = 8000
    pd.options.display.max_rows = 8000
```

In [2]: ┌ df=pd.read_csv('C:/Users/thand/Downloads/house-prices-advanced-regression-techniques.csv')
df.head()

Out[2]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Lvl
0	1	60	RL	65.0	8450	Pave	NaN	Reg		Lvl
1	2	20	RL	80.0	9600	Pave	NaN	Reg		Lvl
2	3	60	RL	68.0	11250	Pave	NaN	IR1		Lvl
3	4	70	RL	60.0	9550	Pave	NaN	IR1		Lvl
4	5	60	RL	84.0	14260	Pave	NaN	IR1		Lvl

5 rows × 81 columns

In [3]: ┌ #Lets have a peak at the columns
print('There are about {} columns'.format(len(df.columns)))
print(df.columns)

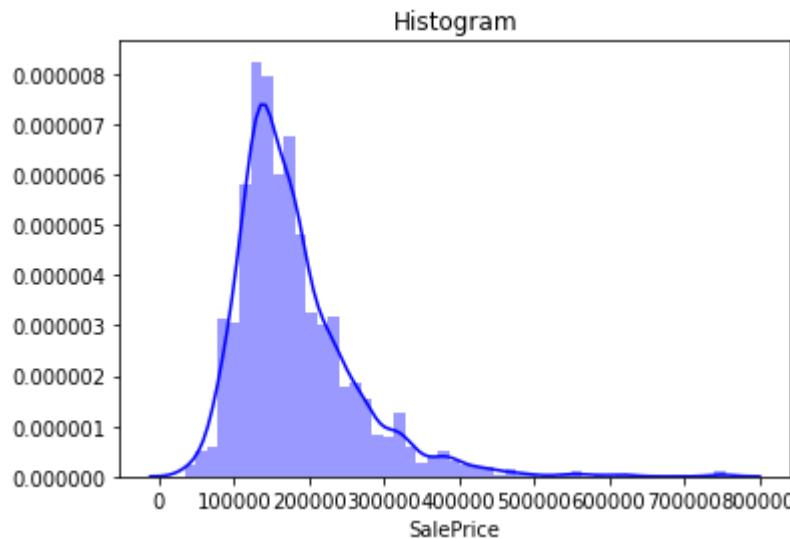
There are about 81 columns

```
Index(['Id', 'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street',
       'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig',
       'LandSlope', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType',
       'HouseStyle', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd',
       'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType',
       'MasVnrArea', 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual',
       'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1',
       'BsmtFinType2', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating',
       'HeatingQC', 'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF',
       'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath',
       'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual',
       'TotRmsAbvGrd', 'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType',
       'GarageYrBlt', 'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual',
       'GarageCond', 'PavedDrive', 'WoodDeckSF', 'OpenPorchSF',
       'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'PoolQC',
       'Fence', 'MiscFeature', 'MiscVal', 'MoSold', 'YrSold', 'SaleType',
       'SaleCondition', 'SalePrice'],
      dtype='object')
```

In [15]: ┌ #Lets take a look at the dependent variable
continuousVariable(df, 'SalePrice')

Skewness:1.8828757597682129

Kurtosis:6.536281860064529



The Dependent variable, SalePrice is Left skewed, Non Normal and Leptokurtic. Needs transformations for linear regression

In [16]: ┌ # Variables by data types

#In this section, we will make the following lists

#num_cols: list of all numerical variables

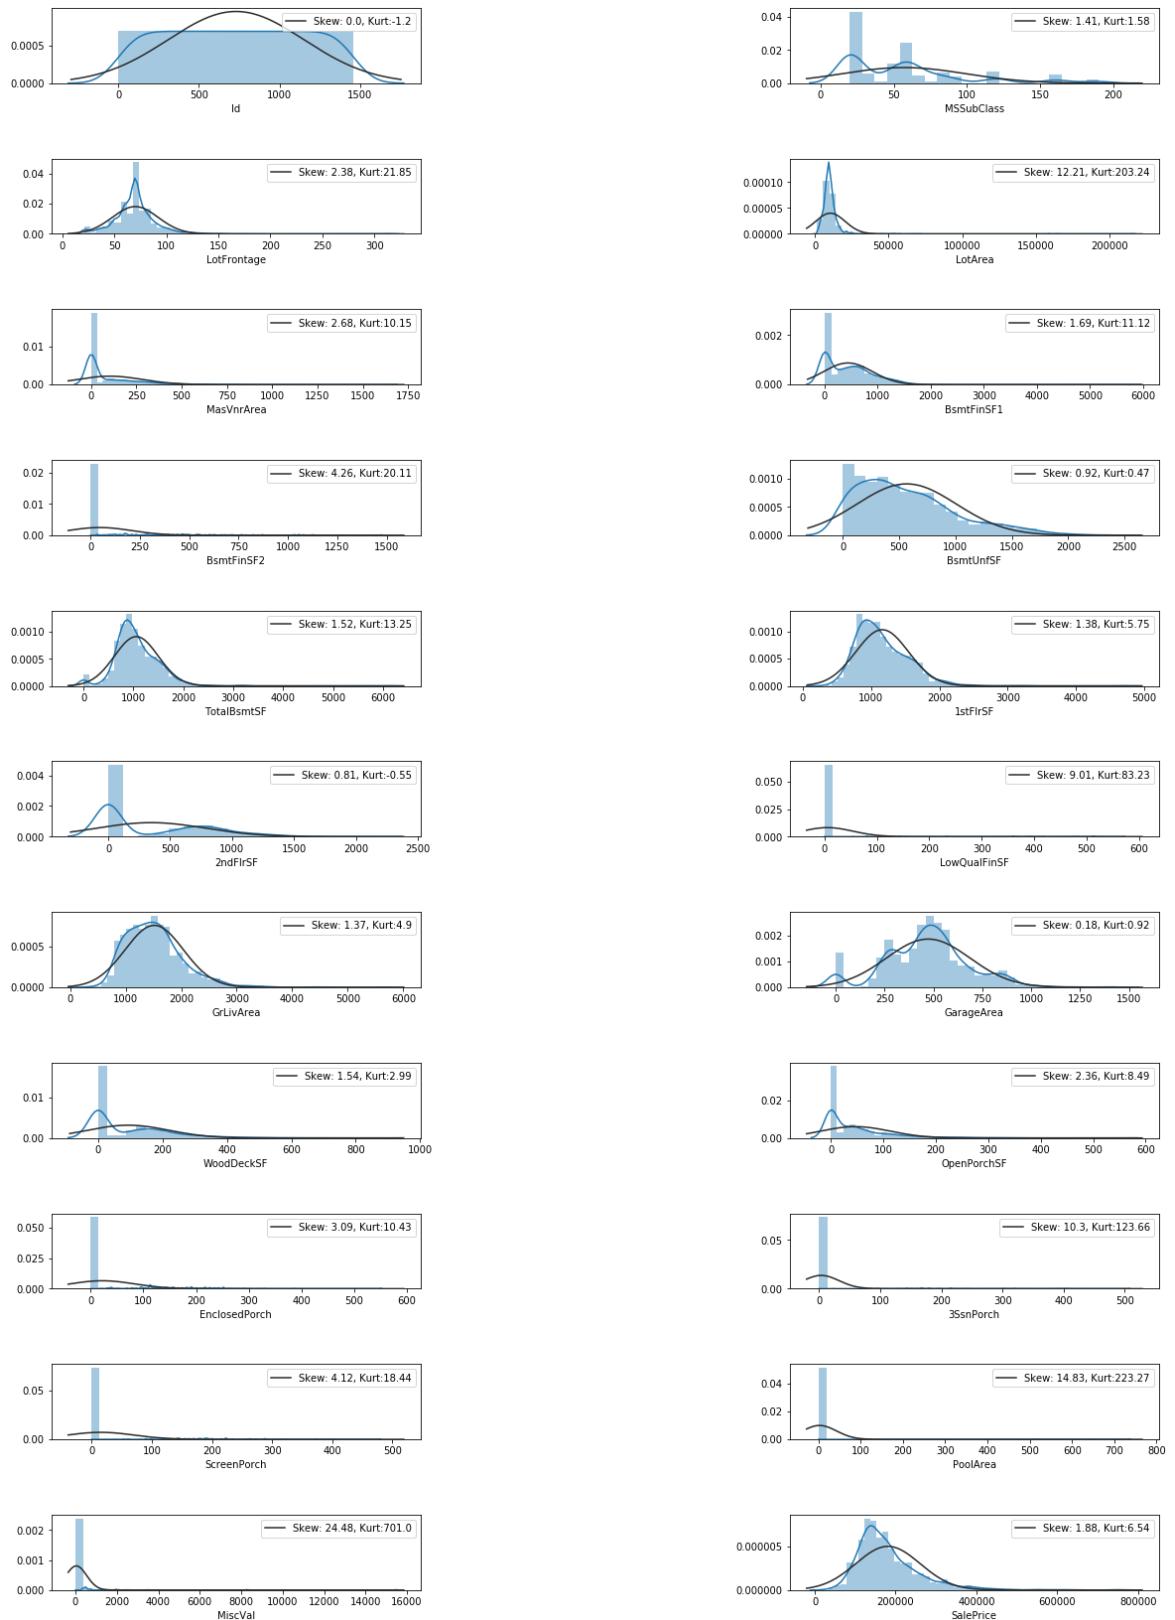
#cat_cols: list of all categorical variables

#ordinal_cols: list of all mean encoded ordinal variables

```
#discrete ratio/interval to categorical features to get more info on them
numericToCategory(df,['MoSold','YrSold','TotRmsAbvGrd','OverallQual','OverallCond'])
numerics = ['int16', 'int32', 'int64', 'float16', 'float32', 'float64']
num_cols= df.select_dtypes(include=numerics).columns #List of all numerical
#Nominal/ordinal
```

In [17]: # Numerical variables distribution

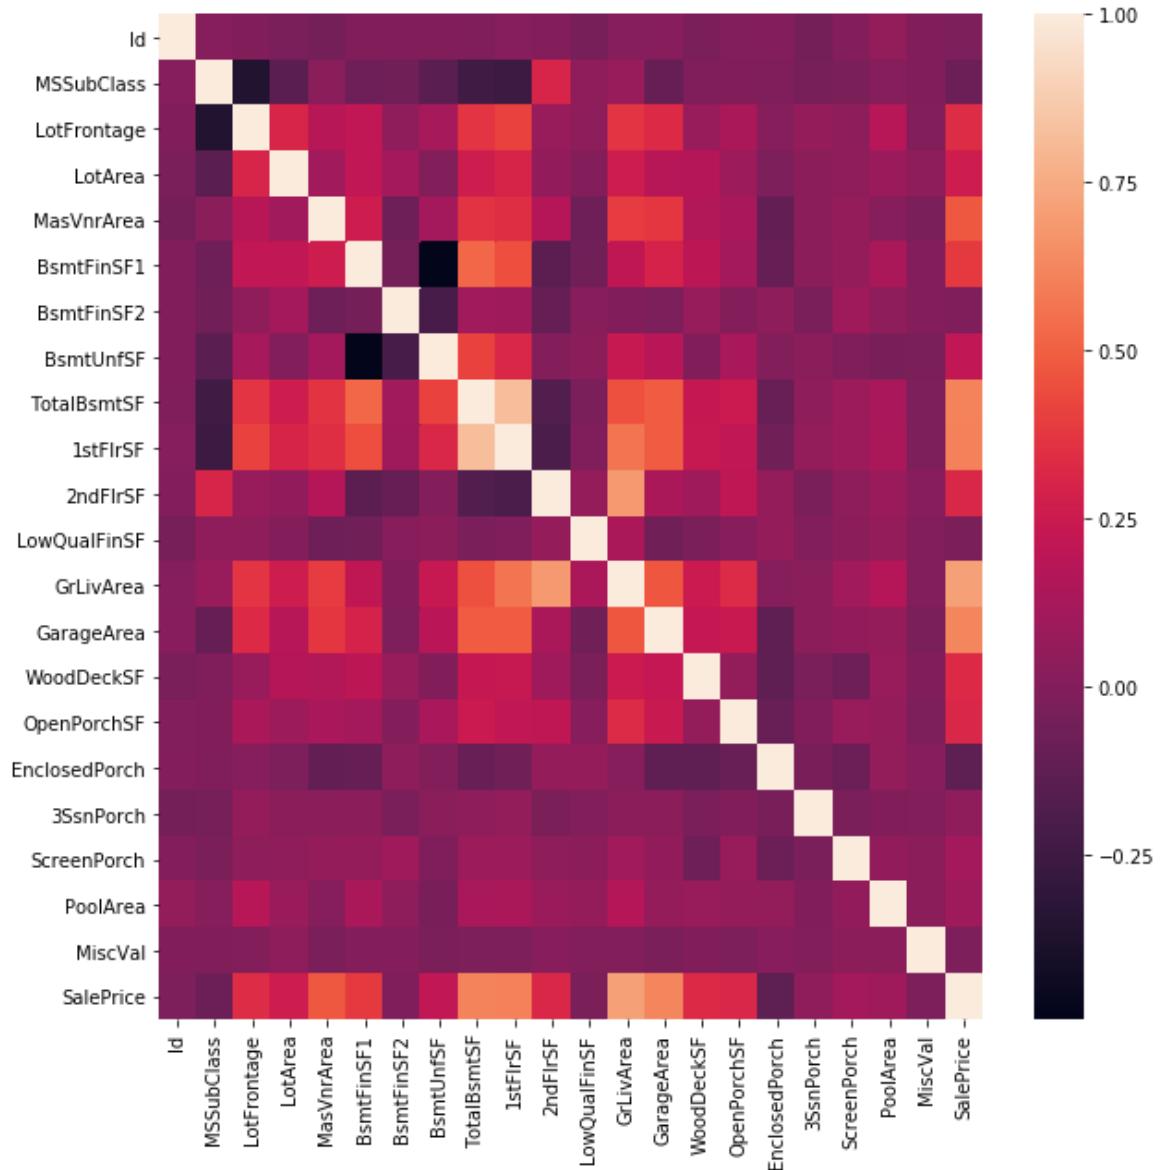
```
NumericalEDA('Histogram',22,2,df,df.select_dtypes(include=numerics).columns,)
```



Most of the variables are not Normal, heavily skewed. Needs processing.

In [19]: ┌ # Lets Look at highly correlated Numerical variables

```
NumericalEDA('Heatmap', 18, 2, df, df.select_dtypes(include=numerics).columns, 'Sa
```





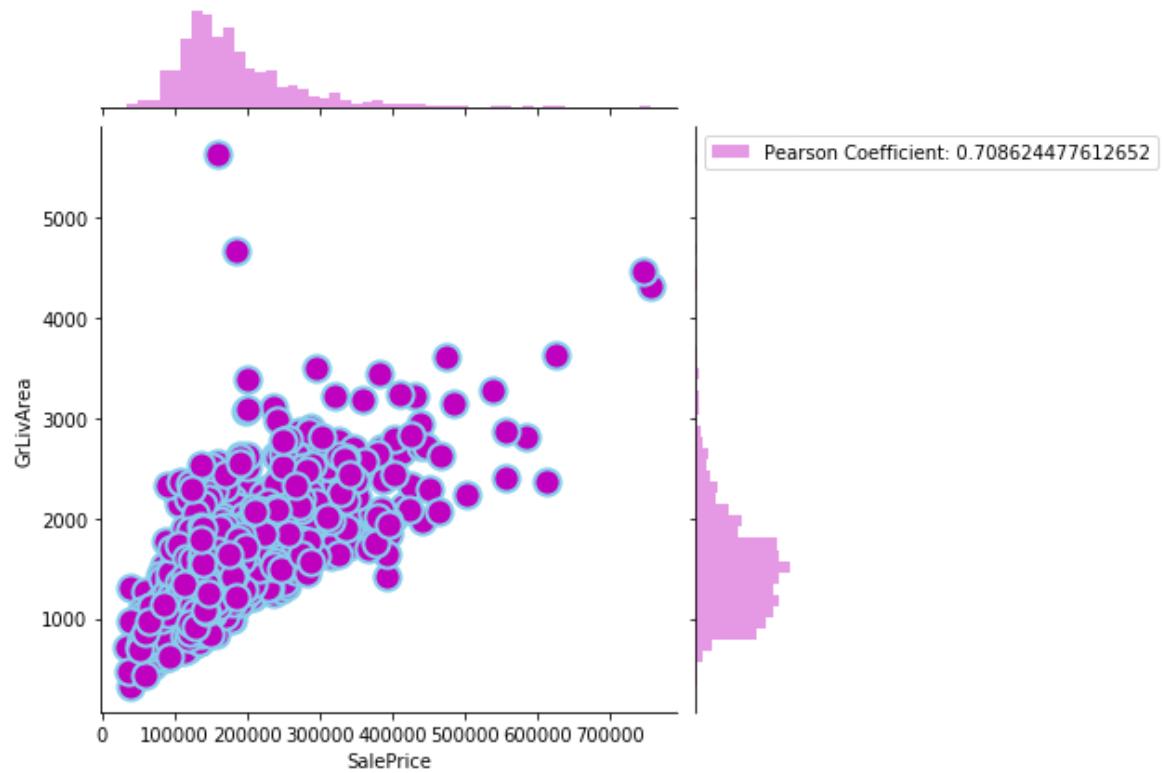
Top 5 highly correlated variables with SalePrice : GrLivArea, GarageArea, TotalBsmtSF, 1stFlrSF, MasVnrArea

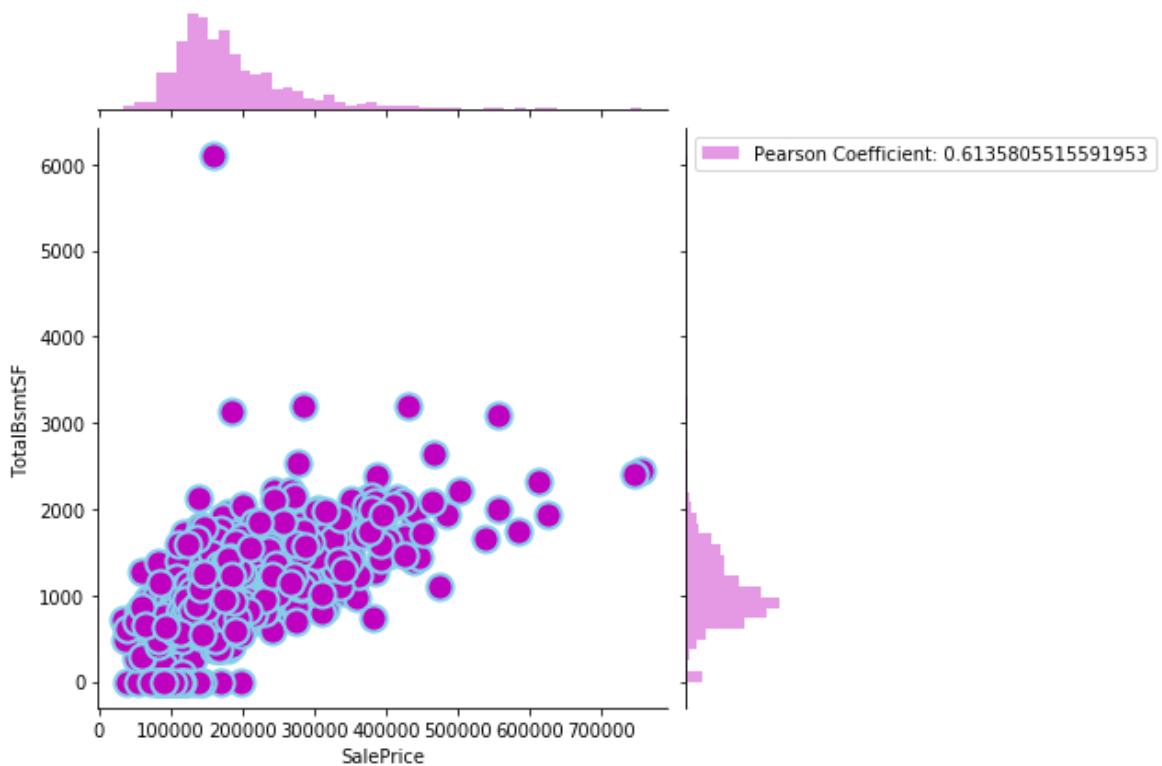
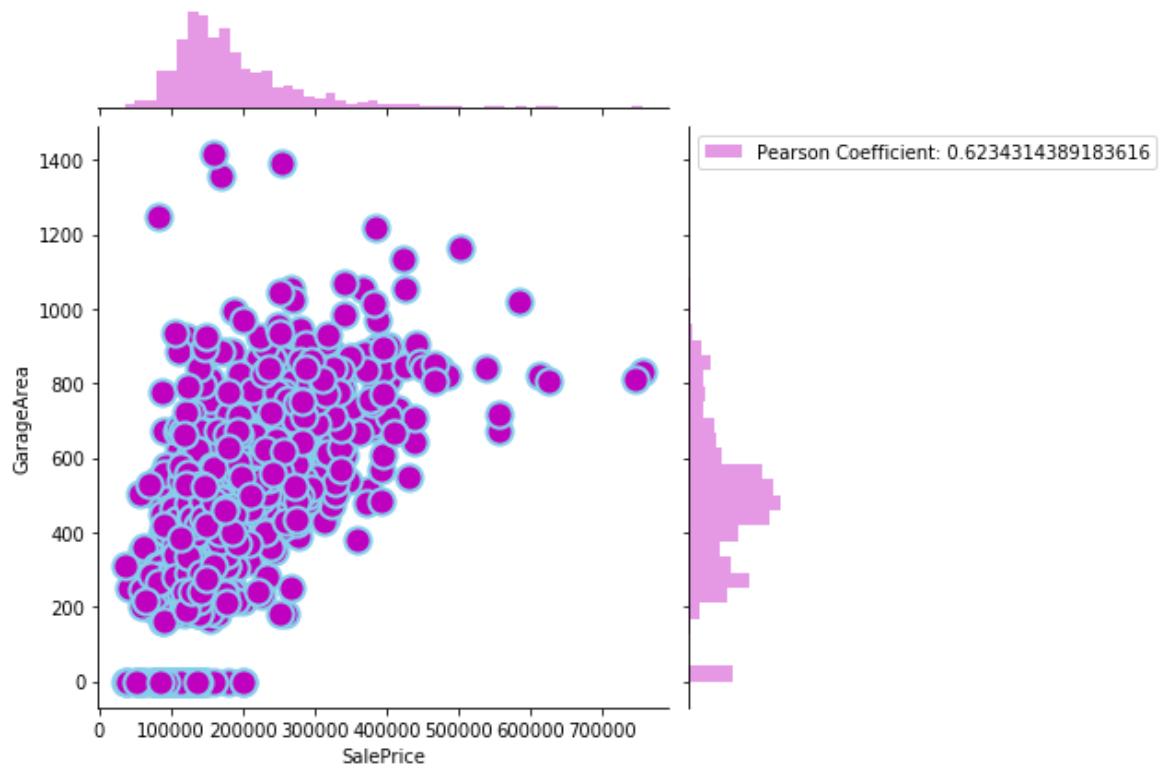
There are Many correlated independent variables as well 'GrLivArea'--> 2ndFlrSF, 1ndFlrSF
 'TotalBsmtSF'-->BsmtFinSF1, 1stFlrSF

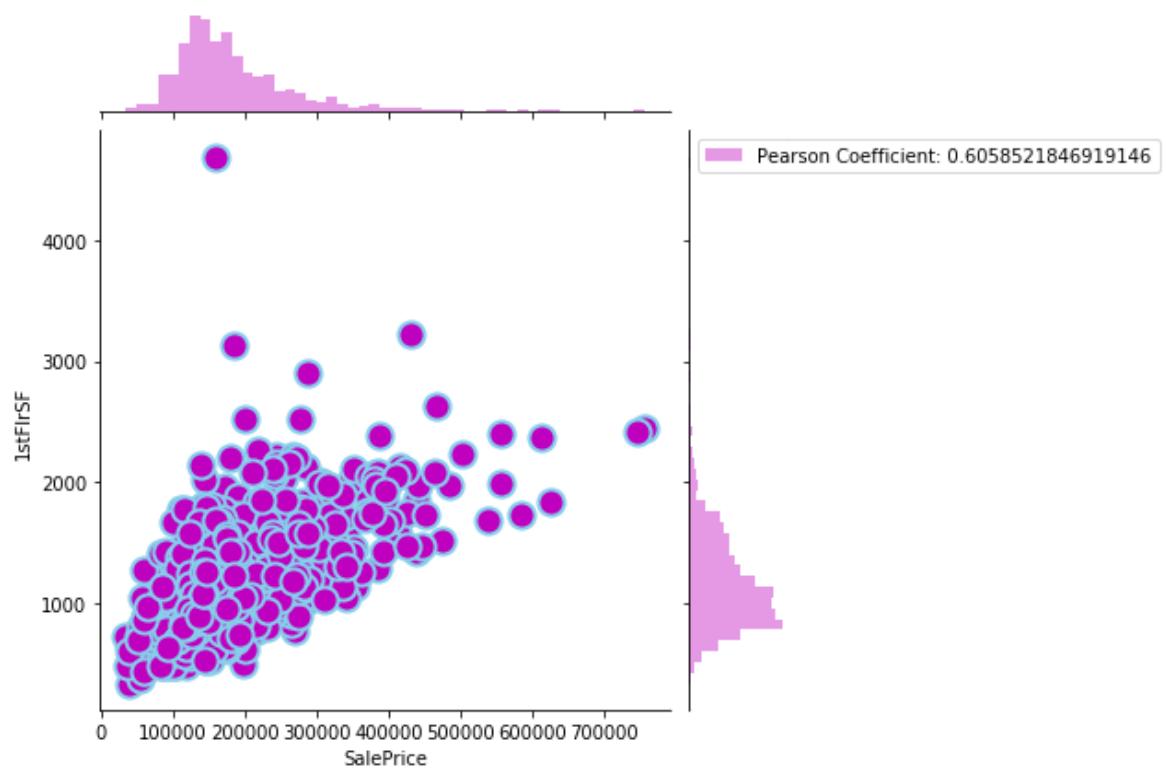
Processing is need while modelling Linear regression

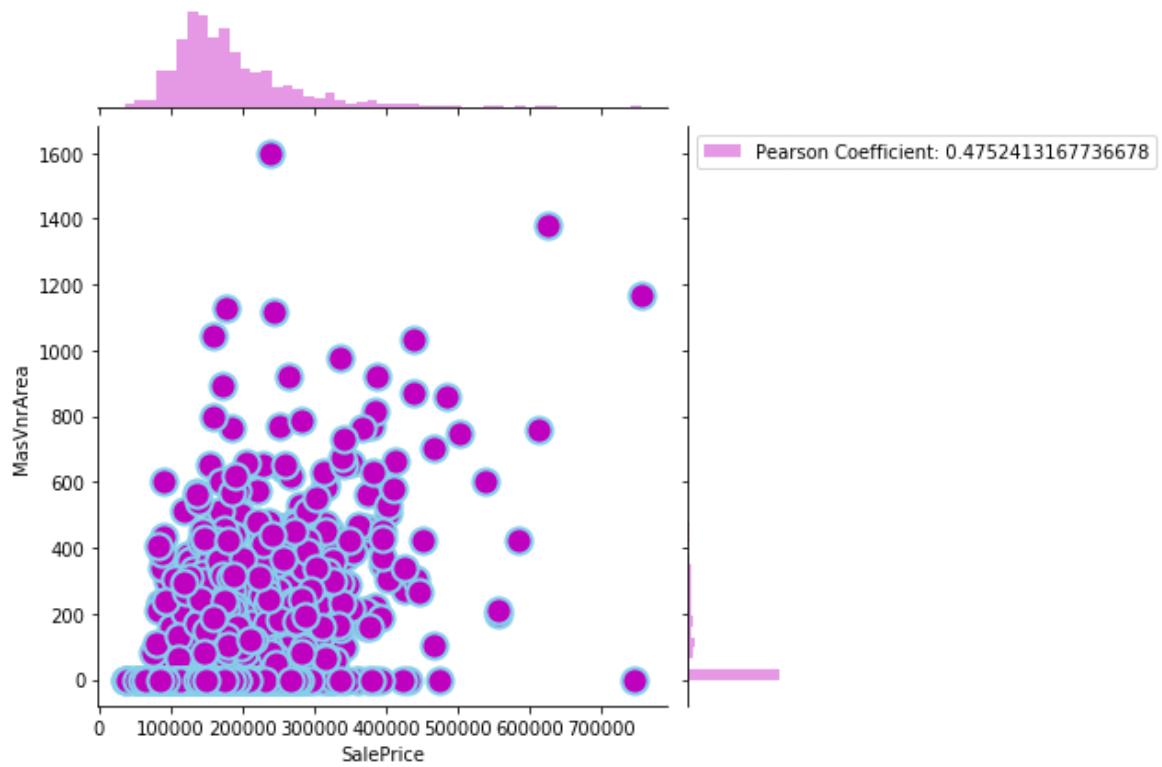
```
In [20]: # Pair plot on the top 5 correlated variables with SalePrice  
NumericalEDA('pair plot',5,1,df,['GrLivArea', 'GarageArea', 'TotalBsmtSF', '1
```

<Figure size 1440x4320 with 0 Axes>







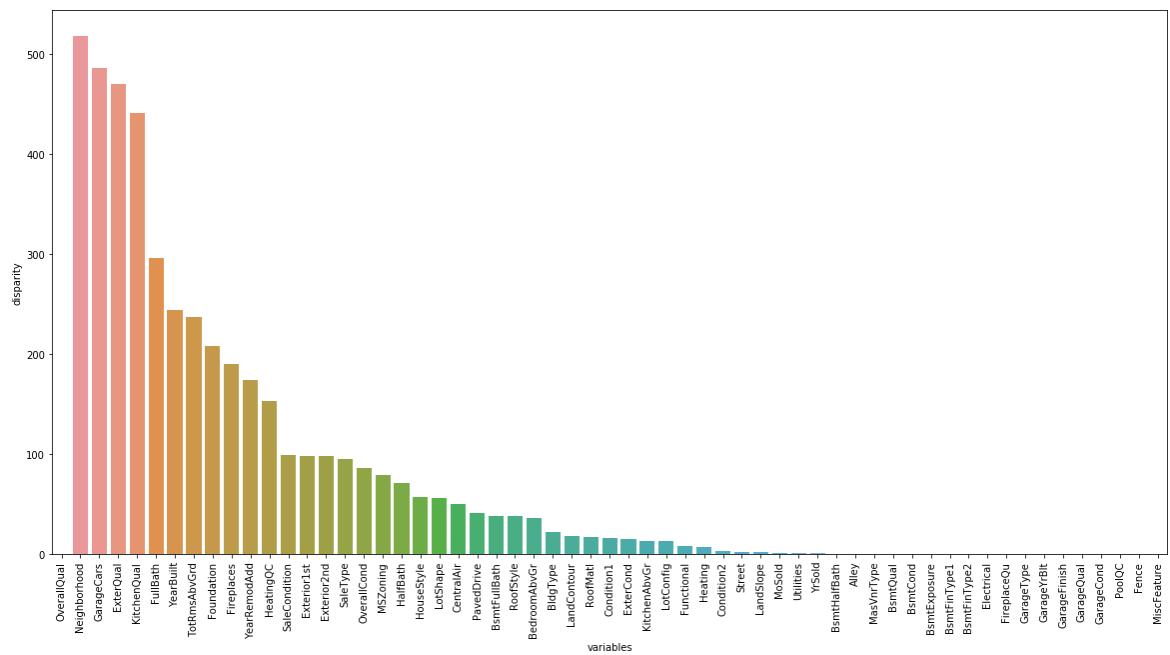


Looking at the scatterplot of the top 5 correlated numerical variables, we observe the following

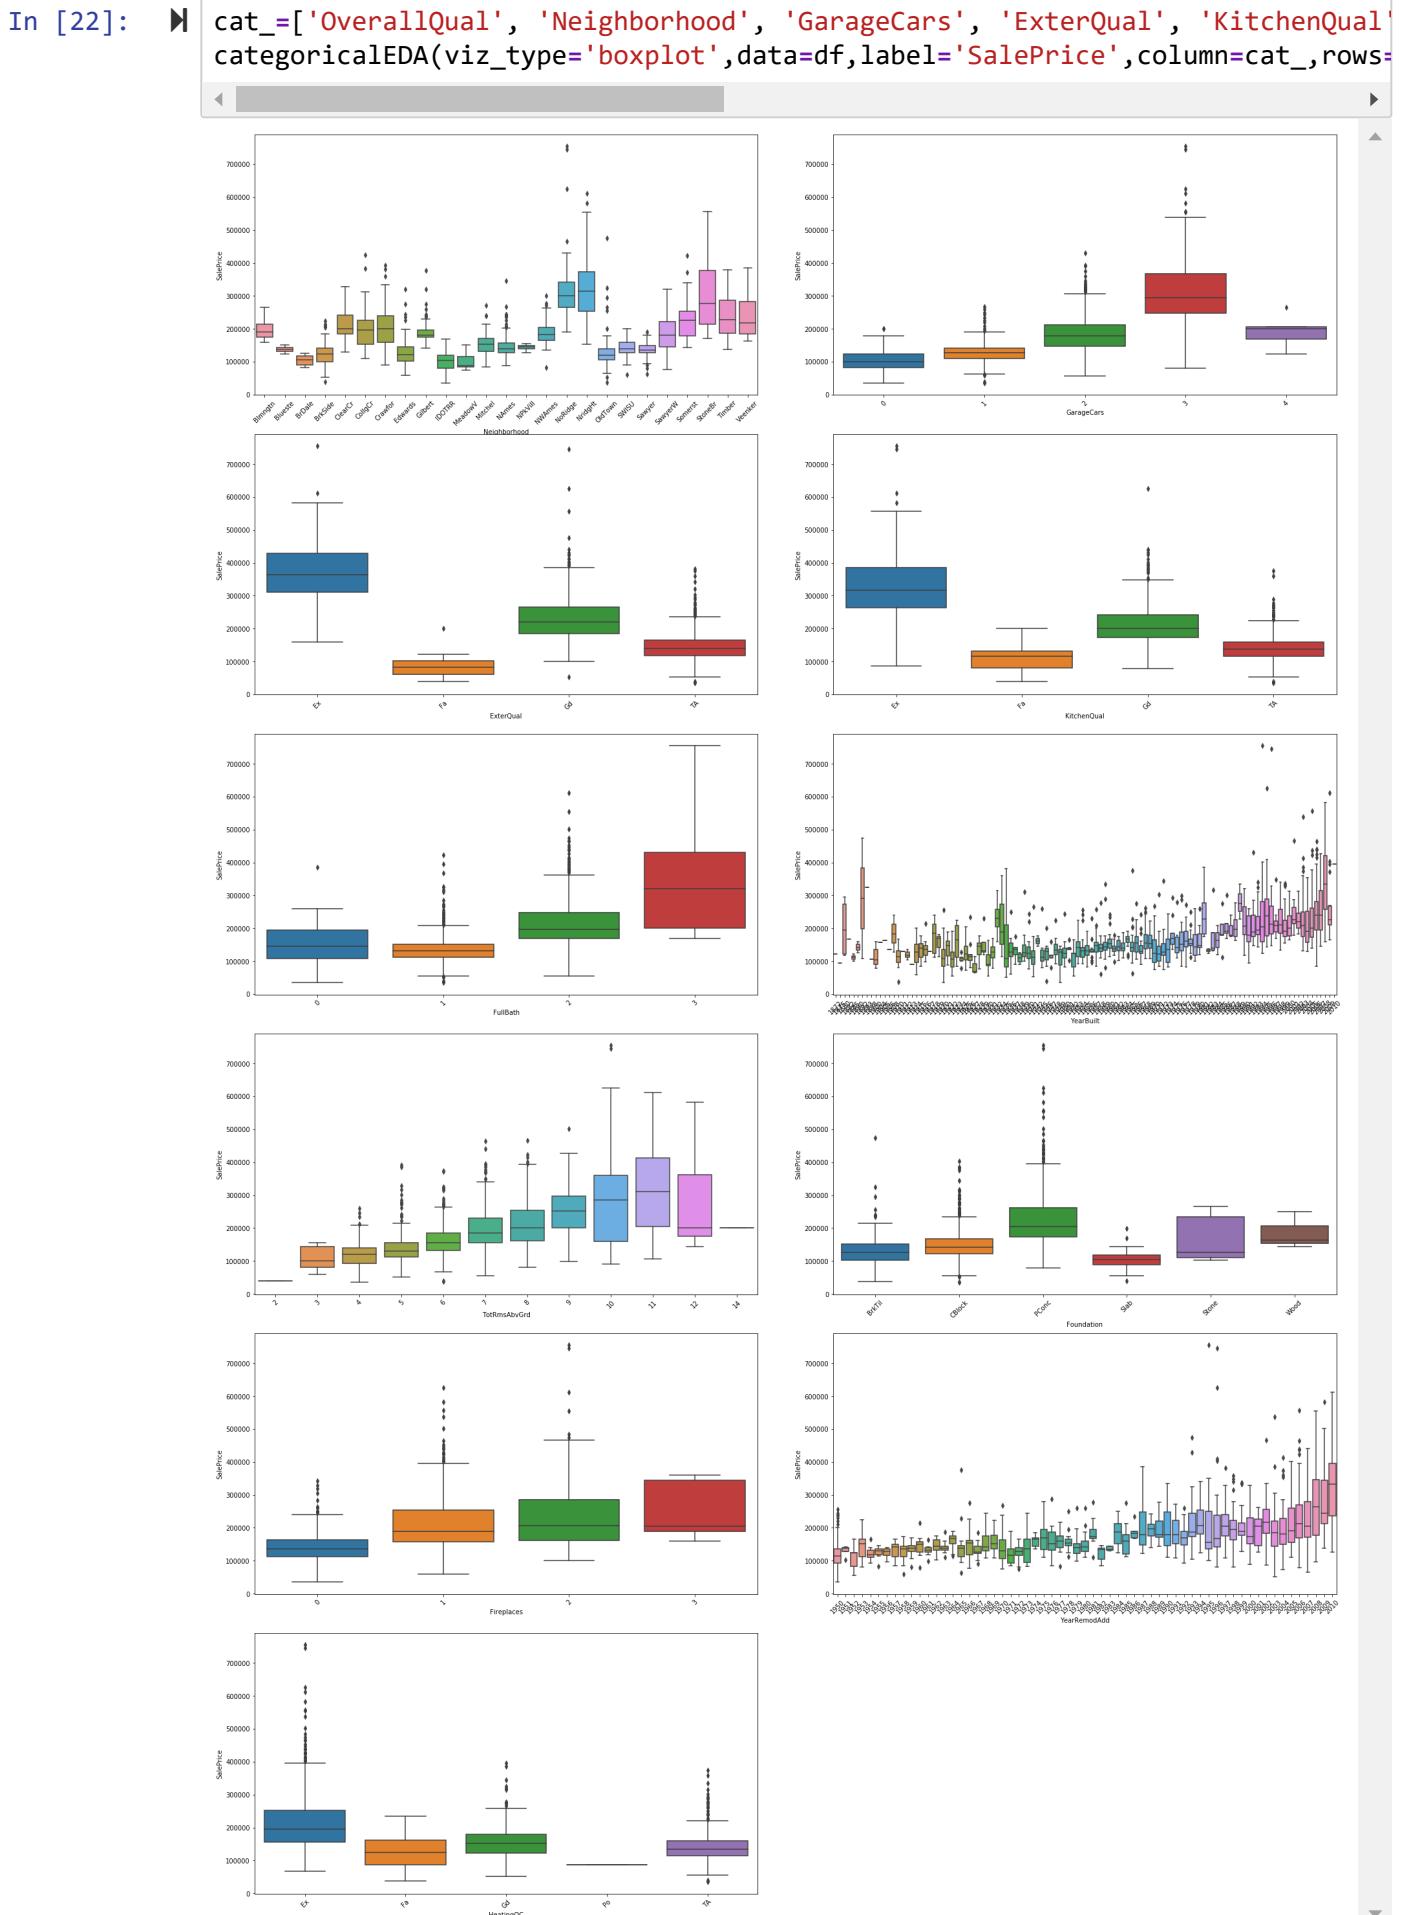
1. There are prominent outliers in the data We need to process these before modelling linear regression

In [21]: # categorical data
1 way anova test

```
categoricalEDA(viz_type='Categorical-VariableImportance', data=df, label='SaleP
```



From the above graph, we can see that OverallQual, Neighborhood, GarageCars, ExterQual, KitchenQual, FullBath, YearBuilt, TotRmsAbvGrd, Foundation, Fireplaces, YearRemodAdd, HeatingQc have significant variability with SalePrice. Lets dig Deep



Houses with the below qualities are significantly priced higher

1. Excellent and Good External Quality
2. Excellent Kitchen Quality
3. 2 or more fullbaths
4. 10+ Rooms
5. Concrete Foundations
6. Excellent heating quality
7. Availability of Fireplace
8. Recently remodelled houses

```
In [5]: def continuousVariable(df=None,col=None):  
    df[col].describe()  
    ax=sns.distplot(df[col],color="Blue")  
    ax.set(xlabel=col, title="Histogram")  
    print("Skewness:{}".format(df[col].skew()))  
    print("Kurtosis:{}".format(df[col].kurt()))
```

```
In [6]: def ordinal_meanencoding(data=None, col=None, label=None, ordinalCols=None):  
    for i in col:  
        a=i+'Ordinal'  
        ordinalCols.append(a)  
        df[a]=df[i].map(dict(zip(df.groupby(i)[label].mean().sort_values(ascending=True),  
                                df[i].unique())))
```

```
In [7]: def NumericalEDA(viz_type=None,rows=None,cols=None,data=None,column=[],label=[])
    if viz_type=='Histogram':
        fig=plt.figure(figsize=(20,60))
        fig.subplots_adjust(hspace=1,wspace=1)
        c=0
        for i in range(1,rows+1):
            if data[column[c]].isnull().any():
                data[column[c]]=data[column[c]].fillna(data[column[c]].mean())
            ax=fig.add_subplot(rows,cols,i)
            ax=sns.distplot(df[column[c]], fit=stats.norm)
            #ax.set(title=column[c])
            plt.legend(["Skew: {}, Kurt:{}".format(round(df[column[c]].skew(),2),round(df[column[c]].kurt(),2))])
            c+=1
        plt.show()

    if viz_type=='Heatmap':
        corrmat=data.corr()
        f,ax=plt.subplots(figsize=(10,10))
        sns.heatmap(corrmat)
        plt.show()

        cols=corrmat.nlargest(10,label)[label].index
        corr=data[cols].corr()
        f,ax=plt.subplots(figsize=(12,9))
        sns.heatmap(corr,cbar=True, annot=True, square=True, fmt=' .2f', annot_kws={"size": 8},cbar_kws={'shrink': 0.5})
        plt.show()

    if viz_type=='joint plot':
        corrmat=data.corr()
        f,ax=plt.subplots(figsize=(10,10))
        sns.heatmap(corrmat)
        plt.show()

        cols=corrmat.nlargest(10,label)[label].index
        f,ax=plt.subplots()
        sns.pairplot(data[cols],size=2.5)
        plt.show()

    if viz_type=='pair plot':
        fig=plt.figure(figsize=(20,60))
        for i in range (1,rows+1):
            #plt.close()
            #fig.add_subplot(rows,cols,i)
            sns.jointplot(x=data[label], y=df[column[i-1]], kind='scatter', s=100)
            plt.legend(["Pearson Coefficient: {}".format(pearsonr(df[label],df[column[i-1]])[0])])
        fig.subplots_adjust(hspace=1,wspace=1)
        plt.show()
```

In [8]: ► `def categoricalEDA(viz_type=None,data=None,column=[],label=None,rows=None,cols=None):`

```

    if viz_type=='boxplot':

        for col in column:
            data[col]=data[col].astype('category')
            if data[col].isnull().any():
                data[col]=data[col].cat.add_categories(['Missing'])
                data[col]=data[col].fillna('Missing')

        f=plt.figure(figsize=(30,50))
        f.subplots_adjust(hspace=0.15, wspace=0.15)
        for i in range(1,rows*2):
            ax=f.add_subplot(rows,cols,i)
            ax=sns.boxplot(x=data[column[i]],y=data[label])
            plt.xticks(rotation=45)
        plt.show()

    if viz_type=='Categorical-VariableImportance':

        anv_df=pd.DataFrame()
        anv_df['variables']=column

        pvals=[]
        catval=[]
        for col in column:
            catval=[]
            for j in df[col].unique():
                catval.append(df.loc[df[col]==j,label])
            #calculate pvalue
            pvals.append(stats.f_oneway(*catval)[1])
        anv_df['pval']=pvals
        anv_df['disparity']=np.log(1/anv_df['pval'].values)
        anv_df.sort_values('disparity',ascending=False,inplace=True)
        fig=plt.figure(figsize=(20,10))
        fig=sns.barplot(x=anv_df['variables'],y=anv_df['disparity'])
        plt.xticks(rotation=90)
        plt.show()

```

In [9]: ► `#missing value`

```

def missing_values(data=None,influential_variable=None,label=None):
    total=data.isnull().sum()
    percent=(data.isnull().sum()/data.isnull().count()).sort_values(ascending=True)
    missing_data=pd.concat([total,percent],axis=1,keys=['total','Percent'],sort_index=True)
    dat=missing_data.loc[missing_data['total']>0,:].reset_index().sort_values('Percent',ascending=False)
    fig=plt.figure(figsize=(20,10))
    ax=sns.barplot(x='index',y='Percent',data=dat)
    ax.set(ylabel="Percent of missing Data")
    ax.set(xlabel="Features")
    ax.set(title="Missing Values (in Percent)")
    plt.xticks(rotation=45)

```

```
In [10]: def numericToCategory(data=None,columns=None):
```

```
    for col in columns:
        data[col]=data[col].astype('category')
```

```
In [11]: def missingValueImputation(data=None,numerical_col=None,categorical_cols=None)
```

```
    number=[]
    cat=[ ]
```

```
    if numerical_col is not None:
```

```
        if numercial_method=='mean_imputation':
            for col in numerical_col:
                data[col]=data[col].fillna(data[col].mean())
```

```
        if numercial_method=='median_imputation':
            for col in numerical_col:
                data[col]=data[col].fillna(data[col].median())
```

```
        if numercial_method=='value_fill':
            for col in numerical_col:
                data[col]=data[col].fillna(numercial_value)
```

```
    for col in number:
        data=data.drop(col,inplace=True)
```

```
    if categorical_cols is not None:
```

```
        if categorical_method=='popular_imputation':
            for col in categorical_cols:
                data[col]=data[col].fillna(data[col].value_counts().index[0])
```

```
        if categorical_method=='value_fill':
            for col in categorical_cols:
                data[col]=data[col].fillna(categorical_value)
```

```
    for col in cat:
        data=data.drop(col,inplace=True)
```

In [12]:

```
n=[]
s=[]

def normalize(data=None, columns=None, std=True):

    if std==True:
        highly_skewed_col=data[columns].columns[df[columns].apply(lambda x: >
f,ax=plt.subplots()
plt.figure(figsize=(10,5))
ax=sns.barplot(x=data[highly_skewed_col].columns,y=data[highly_skewed_col])
ax.xaxis.grid(False)
ax.set(ylabel="Skew")
ax.set(xlabel="Feature names")
ax.set(title="Highly skewed Variables: Before Transformation")
sns.despine(trim=True, left=True)
plt.xticks(rotation=45)
plt.figure(figsize=(10,5))
plt.show()

        for cols in highly_skewed_col:
            power=boxcox_normmax(data[cols] + 1)
            n.append((cols,power))
            data[cols] = stats.boxcox(df[[cols]]+1)[0]

            scalar=StandardScaler().fit(df[columns])
            df[columns] = scalar.transform(df[columns])
            s.append(scalar)

    if std==False:
        highly_skewed_col=data[columns].columns[df[columns].apply(lambda x: >
f,ax=plt.subplots()
plt.figure(figsize=(10,5))
ax=sns.barplot(x=data[highly_skewed_col].columns,y=data[highly_skewed_col])
ax.xaxis.grid(False)
ax.set(ylabel="Skew")
ax.set(xlabel="Feature names")
ax.set(title="Highly skewed Variables: Before Transformation")
sns.despine(trim=True, left=True)
plt.xticks(rotation=45)
plt.show()

        for cols in highly_skewed_col:
            power=boxcox_normmax(data[cols] + 1)
            n.append((cols,power))
            data[cols] = stats.boxcox(df[[cols]]+1)[0]

            ax=sns.barplot(x=data[highly_skewed_col].columns,y=data[highly_skewed_col])
            ax.xaxis.grid(False)
            ax.set(ylabel="Skew")
            ax.set(xlabel="Feature names")
            ax.set(title="Highly skewed Variables: After Transformation")
            sns.despine(trim=True, left=True)
```

```
plt.xticks(rotation=45)
plt.figure(figsize=(10,5))
plt.show()
```

In [13]: ► *# Setup cross validation folds*
kf = KFold(n_splits=12, random_state=42, shuffle=True)

In [14]: ► *# Define error metrics*
def rmsle(y, y_pred):
 return np.sqrt(mean_squared_error(y, y_pred))

def cv_rmse(model, X=None):
 rmse = np.sqrt(-cross_val_score(model, x_train, y_train, scoring="neg_mean_squared_error"))
 return (rmse)

In []: ►